

DRILLS: Debugging RTL Intelligently with Localization from Long Simulation

Specification mining using waveform dumps from long-running FPGA emulation to localize bugs in complex digital designs.

Donggyu Kim & Vighnesh Iyer

Final Presentation: 5/6/2019

Problem Definition

- Higher RTL design productivity (Chisel, HLS) enables more complex digital designs (e.g. BOOM-v2), but complex designs usually contain subtle bugs
- Tricky bugs only manifest after trillions of cycles (e.g. running SPEC2017), as high-level assertion failures, hanging, or termination with errors
 - e.g. pipeline hung, invalid writeback in ROB (see DESSERT [FPL'18])
- Even with a waveform, it is very laborious or even impossible to figure out where and when the bug originated.
- **Problem:** given many error-free traces and one error containing trace, localize the likely bug location by module or line of RTL and find the fine-grained implicit properties which were violated.

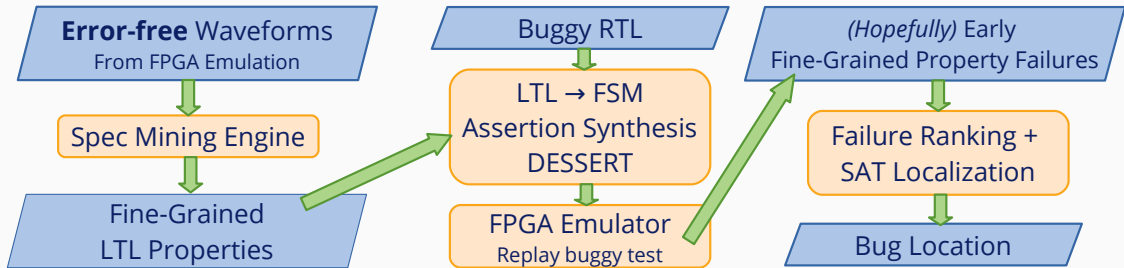
Hypothesis

- *We believe* if a test fails after billions of cycles on a mature RTL design, an assumption the designer made was violated somewhere and at some time
 - The final test failure could have sprung from latent state that was corrupted billions of cycles ago
- The designer's *assumptions* can be extracted from waveforms with 'normal' activity using specification mining
- We can add these mined specifications (as assertions) to the design and replay the failing test to catch a faulty assumption **earlier and with greater locality**

Goals/Approach

- We want a specification mining system that can address:
 - Automated suite of regression assertions
 - Early bug detection and localization on long-running tests
 - Starting point for formal specs of a design
- We aim to first replicate Li's spec mining engine and templates, and apply it to `riscv-mini` on short ISA tests to detect user-inserted bugs
- Next, we will apply the spec mining engine to `rocket-chip` and try it on longer traces

Proposed Approach



LTL (Linear Temporal Logic)

- LTL is a logic for defining *properties* over *traces* of *atomic propositions*
- In the digital design context:
 - A trace is a signal (registers/wires) sampled at rising clock edges; traces are of finite length (not true for LTL in general)
 - Atomic propositions (AP) are functions over signal(s) that evaluate to boolean values

Trace of signal $a = \tau_a = \{0, 1, 1, 1, 0\}$

Trace of signal $b = \tau_b = \{0, 100, 200, 300, 0\}$

AP $x = f(a) = (a == 1)$, AP $y = f(b) = (b \geq 200)$

Trace of AP $x = \tau_x = \{0, 1, 1, 1, 0\}$

Trace of AP $y = \tau_y = \{0, 0, 1, 1, 0\}$

LTL Logical Operators

- An LTL formula ϕ is defined over a trace of tuples of APs.
 - Consider the last example: $\tau = \{(), (x), (x, y), (x, y), ()\}$
- An AP on its own is a valid LTL formula
 - For example $\phi = x$. This formula is satisfied if the first element of the trace contains the AP.
- APs can be composed with simple logic operators ($\neg, \wedge, \vee, \rightarrow$) to form another valid LTL formula
 - For example $\phi = (\neg x \vee y)$. This formula is satisfied if the first element of the trace contains the AP.

LTL Temporal Operators

- There are 4 (commonly used) LTL temporal operators (let p, q be APs or LTL formulas):
 - **G** p (*globally*, p must hold for the entire trace)
 $\{(p), (p), (p, q), (p), (p)\}$
 - **F** p (*eventually*, p must eventually hold at some point)
 $\{(), (), (q), (q, p), (p)\}$
 - **X** p (*next*, p must hold in the next timestep)
 $\{(), (p), (p), (q), ()\}$
 - p **U** q (*until*, p must hold until q holds)
 $\{(p), (p), (p, q), (q), ()\}$
- Common combinations include **GF** p (p holds infinitely often) and **FG** p (once p holds, it holds forever)

Hardware Idioms in LTL

Many common temporal RTL patterns are expressible in LTL.

- There should eventually be a response after a request is dispatched
 $\mathbf{G}(\text{req} \rightarrow \mathbf{XF} \text{resp})$
- The memory bus should respond in 2 cycles
 $\mathbf{G}(\text{req} \rightarrow \mathbf{XX} \text{resp})$
- IrrevocableIO should keep valid high until ready has been asserted
 $\mathbf{G}(\text{valid} \rightarrow \mathbf{X}(\text{valid} \mathbf{U} \text{ready}))$
- After a transaction, the slave should be ready again within 2 cycles
 $\mathbf{G}((\text{valid} \wedge \text{ready}) \rightarrow (\mathbf{X} \text{ready} \vee \mathbf{XX} \text{ready}))$

LTL Templates

There are several LTL formulas that can be 'templated' to be filled in with concrete signals from a design.

- Alternating: $a \mathbf{A} b$
Example: $\{(a), (a), (b), (a), (b)\}$
- Until: $\mathbf{G}(a \rightarrow \mathbf{X}(a \mathbf{U} b))$
- Next: $\mathbf{G}(a \rightarrow \mathbf{X} b)$
- Eventual: $\mathbf{G}(a \rightarrow \mathbf{X} \mathbf{F} b)$

These property templates are from Wenchao Li's thesis. We implement these templates in the spec mining engine.

Delta Traces

- Signals in RTL designs aren't APs on their own, so we use delta traces to capture *changes* in a signal from one timestep to the next
 - Trace of signal $a = \tau_a = \{0, 1, 1, 1, 0\}$
 - Trace of signal $b = \tau_b = \{0, 100, 200, 300, 0\}$
 - Delta trace of signal $a = \tau_{\Delta a} = \{0, 1, 0, 0, 1\}$
 - Delta trace of signal $b = \tau_{\Delta b} = \{0, 1, 1, 1, 1\}$
- VCD (value change dump) files already store delta traces in compressed format
- Custom user-defined events can be used to construct AP traces (such as AP $f(r, v) = r \wedge v$ for ready/valid interfaces)

Specification Mining

- We use a tool (SPOT) to construct a Buchi-automaton that acts like an LTL property monitor




Figure 1: Monitor for $\mathbf{G}(a \rightarrow \mathbf{X} \mathbf{F} b)$

- We use a naive algorithm ($\mathcal{O}(n!)$) to plug in delta traces to the template and check whether the property holds.
 - We use the sparsity of delta traces to perform spec mining efficiently
 - We restrict the signal permutations under consideration to be within 1 module

As of 2 Weeks Ago

- We have a spec miner in Python that can evaluate the 'alternating' and 'next' templates using VCD files as input
- It works on chisel-examples and riscv-mini (and rocket-chip if only 1 clock is considered)


Here's an example property mined from riscv-mini

```
Next: [['TOP.Tile.core.dpath.csr_io_stall',  
  ↳ 'TOP.Tile.core.dpath.stall',  
  ↳ 'TOP.Tile.core.dpath.csr.io_stall'],  
  ↳ ['TOP.Tile.icache.is_idle']]
```



As of 2 Weeks Ago: Bugs Everywhere!

```
Next: [['TOP.Tile.core.dpath.csr_io_stall',  
  ↳ 'TOP.Tile.core.dpath.stall',  
  ↳ 'TOP.Tile.core.dpath.csr.io_stall'],  
  ↳ ['TOP.Tile.icache.is_idle']]
```



Look closely, the waveform doesn't match the property template!

Current Status

- Support for all 4 property templates and some tests
- Support for module-level mining and deduplication
- Merging mined properties from several VCDs
- Checking property sets against VCDs
- Notions of falsifiability, falsification, and support
- Automated mine and check flow for `riscv-mini`

Mining on a store ISA test

```
python miner.py --start-time 12 --signal-bit-limit 4 rv32ui-p-sw.vcd
```

Top 10 properties:

Until TOP.Tile.core.ctrl.io_A_sel -> TOP.Tile.core.ctrl.io_imm_sel, support: 331

Next TOP.Tile.core.dpath.brCond.eq -> TOP.Tile.core.dpath.brCond.neq, support: 291

Next TOP.Tile.core.dpath.brCond.neq -> TOP.Tile.core.dpath.brCond.eq, support: 291

Eventual TOP.Tile.core.ctrl.io_imm_sel -> TOP.Tile.core.dpath.regFile.io_wen, support:
↪ 248

Next TOP.Tile.core_io_icache_req_valid -> TOP.Tile.icache.is_idle, support: 233

Next TOP.Tile.core_io_icache_req_valid -> TOP.Tile.icache.is_read, support: 233

Until TOP.Tile.core_io_icache_req_valid -> TOP.Tile.icache.state, support: 233

Until TOP.Tile.icache.is_idle -> TOP.Tile.icache.is_read, support: 233

Mining on a jump ISA test

```
python miner.py --start-time 12 --signal-bit-limit 4 rv32ui-p-jal.vcd
Until TOP.Tile.core.dpath.io_ctrl_A_sel -> TOP.Tile.core.dpath_io_ctrl_imm_sel,
↪ support: 37
Until TOP.Tile.core.ctrl.io_wb_sel -> TOP.Tile.core.dpath_io_ctrl_imm_sel, support: 34
Next TOP.Tile.core.dpath_io_ctrl_inst_kill -> TOP.Tile.core.ctrl_io_pc_sel, support: 33
Next TOP.Tile.core.dpath_io_ctrl_inst_kill -> TOP.Tile.core.ctrl.io_wb_sel, support: 33
Next TOP.Tile.core.ctrl_io_pc_sel -> TOP.Tile.core.dpath_io_ctrl_inst_kill, support: 33
Next TOP.Tile.core.ctrl_io_pc_sel -> TOP.Tile.core.ctrl.io_wb_sel, support: 33
Next TOP.Tile.core.dpath.brCond.eq -> TOP.Tile.core.dpath.brCond.neq, support: 33
Next TOP.Tile.core.dpath.brCond.neq -> TOP.Tile.core.dpath.brCond.eq, support: 33
```

Bug Localization

Can typos or copy/paste mistakes be caught?

```
    io.csr_cmd := ctrlSignals(11)
-   io.illegal := ctrlSignals(12)
+   io.illegal := ctrlSignals(11)
```

This caused a test to hang, so let's check if any mined specs were violated:

```
python checker.py --start-time 12 --signal-bit-limit 4 rv32ui-p-sub.vcd
↪ riscv_mini.props
ERROR on property Until TOP.Tile.core.dpath.csr.io_illegal ->
↪ TOP.Tile.icache.io_cpu_req_valid
ERROR on property Until TOP.Tile.core.dpath.csr.io_illegal ->
↪ TOP.Tile.icache.io_cpu_resp_valid
ERROR on property Until TOP.Tile.core.dpath.csr.io_illegal ->
↪ TOP.Tile.core.ctrl.io_A_sel
```

Challenges

- Many of the specs are meaningless or require very specific behaviors to falsify
- Some of the specs can be statically inferred (like shift registers)
- These spec templates aren't sufficient to catch some typo bugs
- The VCD parser is too slow to run on longer traces like the RISC-V benchmarks

Current Status and Future Work

- We built a spec mining engine to replicate the prior work
- We applied the engine to `riscv-mini` and demonstrated its strengths and weaknesses
- There's a lot of future work in refining the spec templates and in software engineering
- The ultimate goal of applying this work to debug BOOM is a while away